
```
In [9]: #
import pandas as pd
import numpy as np
```

```
In [10]: # 1.
names = ['Sample code number', 'Clump Thickness', 'Uniformity of Cell Siz
        'Uniformity of Cell Shape', 'Marginal Adhesion', 'Single Epithel
        'Bare Nuclei', 'Bland Chromatin', 'Normal Nucleoli', 'Mitoses',

data_path = "data/breast-cancer-wisconsin.data"
data = pd.read_csv(data_path, names=names)

# 2
# 2.1
data = data.replace(to_replace="?", value=np.nan)
data = data.dropna()

# 2.2
data.iloc[:, 1:] = data.iloc[:, 1:].apply(pd.to_numeric, errors='coerce')
data = data.dropna()

x = data.iloc[:, 1:10].values.astype(np.float64)
y = data["Class"].values.astype(int)
y = np.where(y == 4, 1, 0)
```

```
In [11]: # 2.3
def train_test_split_manual(X, y, test_size=0.25, random_state=2025):
    """
    if random_state is not None:
        np.random.seed(random_state)

    n_samples = X.shape[0]
    n_test = int(n_samples * test_size)

    #
    indices = np.random.permutation(n_samples)
    test_indices = indices[:n_test]
    train_indices = indices[n_test:]

    return X[train_indices], X[test_indices], y[train_indices], y[test_in
```

```
In [12]: # 2.4
class StandardScaler:
    # TODO
    #
    #         : z = (x - mean) / std

    def __init__(self):
        self.mean_ = None
        self.std_ = None

    def fit(self, X):
        """
```

```

self.mean_ = np.mean(X, axis=0)
self.std_ = np.std(X, axis=0)

def transform(self, X):
    """
    return (X - self.mean_) / self.std_

def fit_transform(self, X):
    """
    self.fit(X)
    return self.transform(X)

```

$$h_{\theta}(x) = \frac{1}{1+e^{-\theta x}}$$

- x
- θ
- $h_{\theta}(x)$

$$\text{class} = \begin{cases} 1, & \text{if } h_{\theta}(x) \geq 0.5 \\ 0, & \text{if } h_{\theta}(x) < 0.5 \end{cases}$$

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right]$$

$$\theta := \theta - \alpha \frac{\partial J(\theta)}{\partial \theta}$$

α

```

In [39]: # 3.
class LogisticRegression:
    """
    :
    learning_rate:
    n_iterations:
    """
    def __init__(self, learning_rate=0.01, n_iterations=1000, regularization=0, lambda_reg=0):
        self.learning_rate = learning_rate
        self.n_iterations = n_iterations
        self.weights = None
        self.bias = None
        self.regularization = regularization
        self.lambda_reg = lambda_reg

    def sigmoid(self, z: np.ndarray) -> np.ndarray:
        # TODO: sigmoid

```

```

    # sigmoid(z) = 1 / (1 + e^(-z))
    return 1 / (1 + np.exp(-z))

def fit(self, X: np.ndarray, y: np.ndarray) -> None:
    # TODO:
    #         BatchGradientDescent BGD

    #         :
    # 1.
    # 2.     n_iterations
    # 3.     :
    #     step1.
    #     step2.
    #     step3.
    n_samples, n_features = X.shape
    self.weights = np.zeros(n_features)
    self.bias = 0

    for _ in range(self.n_iterations):
        linear_pred = np.dot(X, self.weights) + self.bias
        predictions = self.sigmoid(linear_pred)

        d_weights = np.dot(X.T, (predictions - y)) / n_samples
        d_bias = np.sum(predictions - y) / n_samples

        if self.regularization == 'L1':
            d_weights += self.lambda_reg * np.sign(self.weights)
        elif self.regularization == 'L2':
            d_weights += 2 * self.lambda_reg * self.weights

        self.weights -= self.learning_rate * d_weights
        self.bias -= self.learning_rate * d_bias

def predict(self, X: np.ndarray, threshold: float = 0.5) -> np.ndarray:
    # TODO:
    #         :
    # 1.
    # 2.     sigmoid
    # 3.
    linear_pred = np.dot(X, self.weights) + self.bias
    predictions = self.sigmoid(linear_pred)
    return np.where(predictions >= threshold, 1, 0)

def predict_proba(self, X: np.ndarray) -> np.ndarray:
    linear_pred = np.dot(X, self.weights) + self.bias
    return self.sigmoid(linear_pred)

```

```

In [14]: def get_metrics(y_true, y_pred):
    """

    """

    def recall_score(y_true, y_pred):
        """

        = TP / (TP + FN)
        """
        TP = np.sum((y_true == 1) & (y_pred == 1))
        FN = np.sum((y_true == 1) & (y_pred == 0))
        return TP / (TP + FN) if TP + FN > 0 else 0

```

```

def precision_score(y_true, y_pred):
    """
    = TP / (TP + FP)"""
    TP = np.sum((y_true == 1) & (y_pred == 1))
    FP = np.sum((y_true == 0) & (y_pred == 1))
    return TP / (TP + FP) if TP + FP > 0 else 0

def accuracy_score(y_true, y_pred):
    """
    = (TP + TN) / """
    return np.mean(y_true == y_pred)

def confusion_matrix(y_true, y_pred):
    """
    """
    TN = np.sum((y_true == 0) & (y_pred == 0))
    FP = np.sum((y_true == 0) & (y_pred == 1))
    FN = np.sum((y_true == 1) & (y_pred == 0))
    TP = np.sum((y_true == 1) & (y_pred == 1))

    return np.array([[TN, FP], [FN, TP]])

recall = recall_score(y_true, y_pred)
precision = precision_score(y_true, y_pred)
accuracy = accuracy_score(y_true, y_pred)
cm = confusion_matrix(y_true, y_pred)
return recall, precision, accuracy, cm

```

L1 L2

```

In [31]: def compare_regularization_effects():
    """
    """
    print("=" * 80)
    print(" ")
    print("=" * 80)

    #
    X_train, X_test, y_train, y_test = train_test_split_manual(x, y, test

    scaler = StandardScaler()
    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)

    #
    models = {
        ': LogisticRegression(learning_rate=0.01, n_iterations=10
        'L1': LogisticRegression(learning_rate=0.01, n_iterations=10
            regularization='l1', lamb
        'L2': LogisticRegression(learning_rate=0.01, n_iterations=10
            regularization='l2', lamb
    }

    results = {}

    for name, model in models.items():
        print(f"\n {name} ...")
        model.fit(X_train, y_train)

        #
        y_pred = model.predict(X_test)

```

```

recall, precision, accuracy, cm = get_metrics(y_test, y_pred)

results[name] = {
    'recall': recall,
    'precision': precision,
    'accuracy': accuracy,
    'weights_norm': np.linalg.norm(model.weights),
    'weights_sparsity': np.sum(np.abs(model.weights) < 1e-6) / len(r
}

print(f"{name}      :")
print(f"      : {recall:.4f}")
print(f"      : {precision:.4f}")
print(f"      : {accuracy:.4f}")
print(f"      : {np.linalg.norm(model.weights):.4f}")
print(f"      : {np.sum(np.abs(model.weights) < 1e-6) / len(r

return results

```

```

In [36]: def analyze_threshold_effects():
    """ """
    print("\n" + "=" * 80)
    print(" ")
    print("=" * 80)

    #
    X_train, X_test, y_train, y_test = train_test_split_manual(x, y, test

    scaler = StandardScaler()
    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)

    model = LogisticRegression(learning_rate=0.01, n_iterations=1000)
    model.fit(X_train, y_train)

    #
    y_proba = model.predict_proba(X_test)

    #
    thresholds = np.arange(0.1, 1.0, 0.1)
    threshold_results = []

    print(f"{'<8} {'<10} {'<10} {'<10} {'F1
    print("-" * 50)

    for threshold in thresholds:
        y_pred = (y_proba >= threshold).astype(int)
        recall, precision, accuracy, cm = get_metrics(y_test, y_pred)
        f1_score = 2 * (precision * recall) / (precision + recall) if (pr

        threshold_results.append({
            'threshold': threshold,
            'recall': recall,
            'precision': precision,
            'accuracy': accuracy,
            'f1_score': f1_score
        })

```

```

        print(f"{threshold:<8.1f} {recall:<10.4f} {precision:<10.4f} {acc

#
best_f1_idx = np.argmax([r['f1_score'] for r in threshold_results])
best_threshold = threshold_results[best_f1_idx]

print(f"\n      : {best_threshold['threshold']:.1f}")
print(f"    F1   : {best_threshold['f1_score']:.4f}")

return threshold_results, best_threshold

```

```

In [41]: import matplotlib.pyplot as plt

def plot_threshold_analysis(threshold_results):
    """
    """
    thresholds = [r['threshold'] for r in threshold_results]
    recalls = [r['recall'] for r in threshold_results]
    precisions = [r['precision'] for r in threshold_results]
    f1_scores = [r['f1_score'] for r in threshold_results]

    plt.figure(figsize=(12, 8))

    plt.subplot(2, 2, 1)
    plt.plot(thresholds, recalls, 'b-o', label='Recall')
    plt.xlabel('Threshold')
    plt.ylabel('Recall')
    plt.title('Recall vs Threshold')
    plt.grid(True)

    plt.subplot(2, 2, 2)
    plt.plot(thresholds, precisions, 'r-o', label='Precision')
    plt.xlabel('Threshold')
    plt.ylabel('Precision')
    plt.title('Precision vs Threshold')
    plt.grid(True)

    plt.subplot(2, 2, 3)
    plt.plot(thresholds, f1_scores, 'g-o', label='F1-Score')
    plt.xlabel('Threshold')
    plt.ylabel('F1-Score')
    plt.title('F1-Score vs Threshold')
    plt.grid(True)

    plt.subplot(2, 2, 4)
    plt.plot(thresholds, recalls, 'b-o', label='Recall')
    plt.plot(thresholds, precisions, 'r-o', label='Precision')
    plt.plot(thresholds, f1_scores, 'g-o', label='F1-Score')
    plt.xlabel('Threshold')
    plt.ylabel('Metric Value')
    plt.title('All Metrics vs Threshold')
    plt.legend()
    plt.grid(True)

    plt.tight_layout()
    plt.show()

```

```

def plot_regularization_comparison(results):
    """ """
    models = list(results.keys())
    recalls = [results[m]['recall'] for m in models]
    precisions = [results[m]['precision'] for m in models]
    weights_norms = [results[m]['weights_norm'] for m in models]

    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

    #
    x = np.arange(len(models))
    width = 0.35

    ax1.bar(x - width/2, recalls, width, label='Recall', alpha=0.8)
    ax1.bar(x + width/2, precisions, width, label='Precision', alpha=0.8)
    ax1.set_xlabel('Model')
    ax1.set_ylabel('Metric Value')
    ax1.set_title('Regularization Methods Performance Comparison')
    ax1.set_xticks(x)
    ax1.set_xticklabels(models, rotation=45)
    ax1.legend()
    ax1.grid(True, alpha=0.3)

    #
    ax2.bar(models, weights_norms, alpha=0.8, color='orange')
    ax2.set_xlabel('Model')
    ax2.set_ylabel('Weight Norm')
    ax2.set_title('Weight Norm Comparison')
    ax2.tick_params(axis='x', rotation=45)
    ax2.grid(True, alpha=0.3)

    plt.tight_layout()
    plt.show()

```

```

In [8]: if __name__ == '__main__':
        print("=" * 60)
        print("          -          ")
        print("=" * 60)

        #
        X_train, X_test, y_train, y_test = train_test_split_manual(x, y, test

        print(f"\n          :")
        print(f"          : {X_train.shape[0]}")
        print(f"          : {X_test.shape[0]}")
        print(f"          : {X_train.shape[1]}")

        #
        scaler = StandardScaler()
        X_train = scaler.fit_transform(X_train)
        X_test = scaler.transform(X_test)

        print("\n          ...")
        #
        model = LogisticRegression(learning_rate=0.01, n_iterations=1000)
        model.fit(X_train, y_train)

        #
        y_pred = model.predict(X_test)

```

```

#
print("\n" + "=" * 60)
print(" ")
print("=" * 60)

recall, precision, accuracy, cm = get_metrics(y_test, y_pred)

print(f"\n      (Recall):    {recall:.4f} ({recall*100:.2f}%)"
      (Precision): {precision:.4f} ({precision*100:.2f}%)"
      (Accuracy):  {accuracy:.4f} ({accuracy*100:.2f}%)"

print(f"\n      :")
print(f"                                ")
print(f"          {cm[0,0]:4d}          {cm[0,1]:4d}")
print(f"          {cm[1,0]:4d}          {cm[1,1]:4d}")

```

```
=====
```

```
-
```

```
=====
```

```

:
: 513
: 170
: 9
...

```

```
=====
```

```
=====
```

```

(Recall):    0.9831 (98.31%)
(Precision): 0.9667 (96.67%)
(Accuracy):  0.9824 (98.24%)

```

```
:
```

```

      109      2
      1      58

```

In [42]: `if __name__ == '__main__':`

```

#
#
#
print("\n" + "=" * 80)
print(" ")
print("=" * 80)

#
reg_results = compare_regularization_effects()

#
threshold_results, best_threshold = analyze_threshold_effects()

#
plot_threshold_analysis(threshold_results)
plot_regularization_comparison(reg_results)

#
print("\n" + "=" * 80)

```



```
print(" ")
print("=" * 80)

print("\n1.      :")
print("  - L1      ")
print("  - L2      ")
print("  -      ")

print("\n2.      :")
print(f"  -      ")
print(f"  -      ")
print(f"  -      {best_threshold['threshold']:.1f}")
```

```

=====
=====

=====
=====

=====
=====

=====
=====

```

```

      ...
      :
      : 0.9831
      : 0.9667
      : 0.9824
      : 1.5988
      : 0.0000

```

```

L1      L1      ...
L1      :
      : 0.9831
      : 0.9667
      : 0.9824
      : 1.5988
      : 0.0000

```

```

L2      L2      ...
L2      :
      : 0.9831
      : 0.9667
      : 0.9824
      : 1.5988
      : 0.0000

```

```

=====
=====

=====
=====

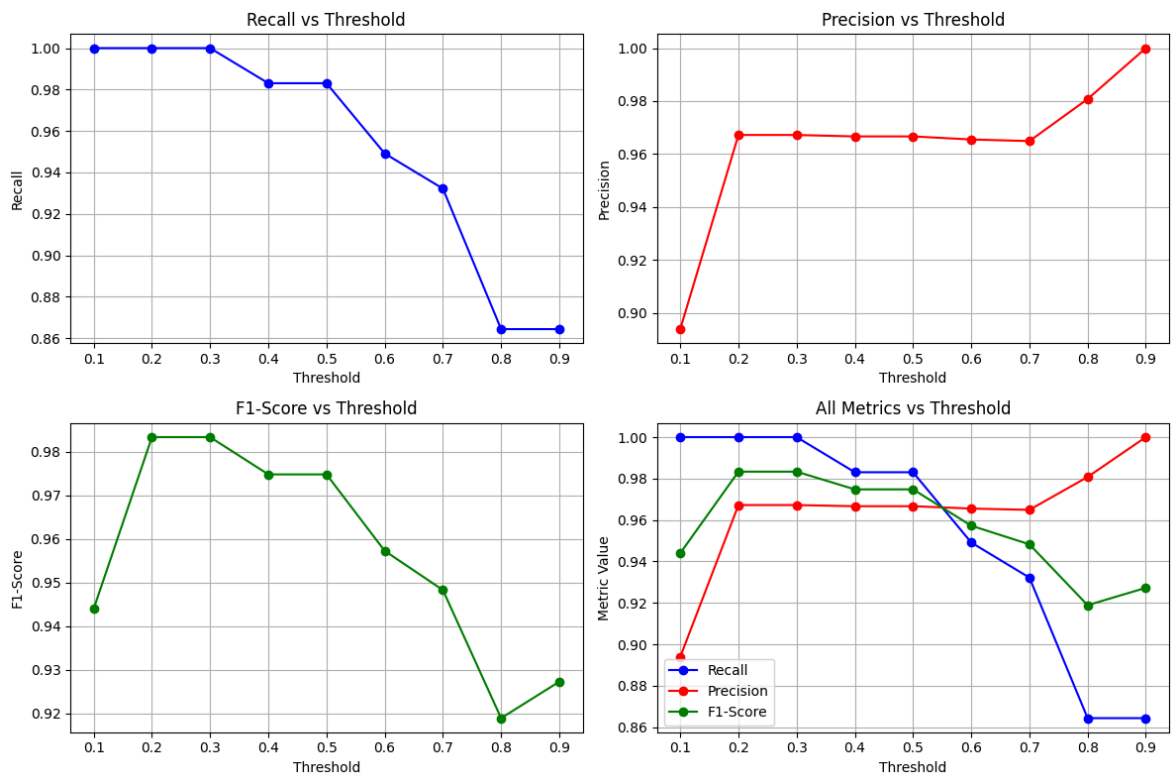
```

				F1
0.1	1.0000	0.8939	0.9588	0.9440
0.2	1.0000	0.9672	0.9882	0.9833
0.3	1.0000	0.9672	0.9882	0.9833
0.4	0.9831	0.9667	0.9824	0.9748
0.5	0.9831	0.9667	0.9824	0.9748
0.6	0.9492	0.9655	0.9706	0.9573
0.7	0.9322	0.9649	0.9647	0.9483
0.8	0.8644	0.9808	0.9471	0.9189
0.9	0.8644	1.0000	0.9529	0.9273

```

      : 0.2
F1    : 0.9833

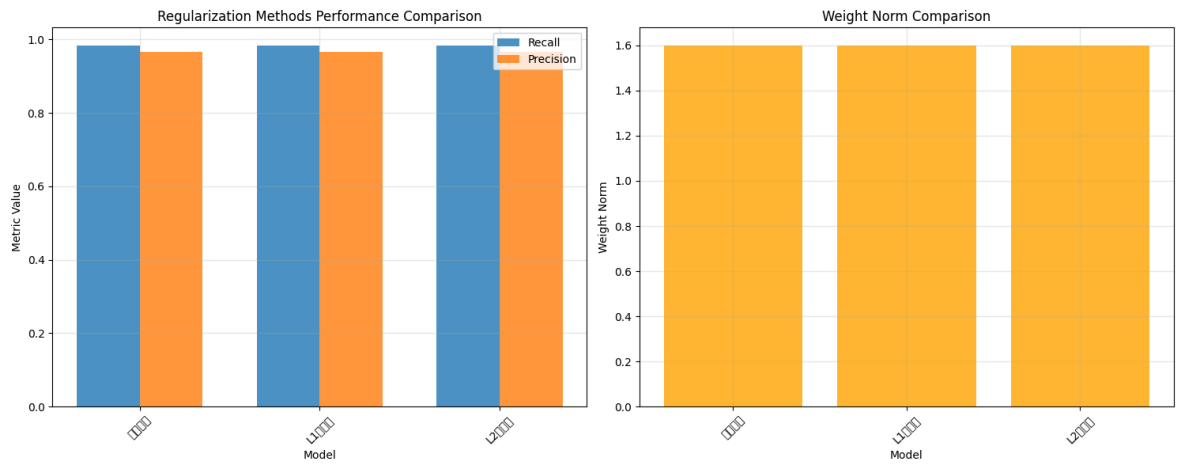
```



```

/var/folders/m7/nm2vlx1n4mg_rg3xtr8gbr640000gp/T/ipykernel_21364/157641164
1.py:77: UserWarning: Glyph 26080 (\N{CJK UNIFIED IDEOGRAPH-65E0}) missing
from font(s) DejaVu Sans.
    plt.tight_layout()
/var/folders/m7/nm2vlx1n4mg_rg3xtr8gbr640000gp/T/ipykernel_21364/157641164
1.py:77: UserWarning: Glyph 27491 (\N{CJK UNIFIED IDEOGRAPH-6B63}) missing
from font(s) DejaVu Sans.
    plt.tight_layout()
/var/folders/m7/nm2vlx1n4mg_rg3xtr8gbr640000gp/T/ipykernel_21364/157641164
1.py:77: UserWarning: Glyph 21017 (\N{CJK UNIFIED IDEOGRAPH-5219}) missing
from font(s) DejaVu Sans.
    plt.tight_layout()
/var/folders/m7/nm2vlx1n4mg_rg3xtr8gbr640000gp/T/ipykernel_21364/157641164
1.py:77: UserWarning: Glyph 21270 (\N{CJK UNIFIED IDEOGRAPH-5316}) missing
from font(s) DejaVu Sans.
    plt.tight_layout()
/Users/Zhuanz/Documents/Senior/PRML/.conda/lib/python3.12/site-packages/IP
ython/core/pylabtools.py:170: UserWarning: Glyph 26080 (\N{CJK UNIFIED IDE
OGRAPH-65E0}) missing from font(s) DejaVu Sans.
    fig.canvas.print_figure(bytes_io, **kw)
/Users/Zhuanz/Documents/Senior/PRML/.conda/lib/python3.12/site-packages/IP
ython/core/pylabtools.py:170: UserWarning: Glyph 27491 (\N{CJK UNIFIED IDE
OGRAPH-6B63}) missing from font(s) DejaVu Sans.
    fig.canvas.print_figure(bytes_io, **kw)
/Users/Zhuanz/Documents/Senior/PRML/.conda/lib/python3.12/site-packages/IP
ython/core/pylabtools.py:170: UserWarning: Glyph 21017 (\N{CJK UNIFIED IDE
OGRAPH-5219}) missing from font(s) DejaVu Sans.
    fig.canvas.print_figure(bytes_io, **kw)
/Users/Zhuanz/Documents/Senior/PRML/.conda/lib/python3.12/site-packages/IP
ython/core/pylabtools.py:170: UserWarning: Glyph 21270 (\N{CJK UNIFIED IDE
OGRAPH-5316}) missing from font(s) DejaVu Sans.
    fig.canvas.print_figure(bytes_io, **kw)

```



1. :

- L1
- L2
-

2. :

-
-
- 0.2

1. sigmoid

- : $\text{sigmoid}(z) \in (0,1)$ 0 1
- :
- : (0, 0.5)
- : $s'(z) = s(z)(1-s(z))$

2. MSE

- - MSE: $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$
 - : $J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$
- - : MSE sigmoid
 - : MSE sigmoid 0
 - :

Recall

$$= TP/(TP+FN)$$

"

"

softmax

1. softmax

- : $\mathbf{x} = [x_1, x_2, \dots, x_n]$
- : $\mathbf{p} = [p_1, p_2, \dots, p_n]$ $\sum_{i=1}^n p_i = 1$

2. softmax

- : softmax
- :
- :

3. softmax

$$\text{softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}, \quad i = 1, \dots, n$$

```
def softmax_stable(x):  
    # softmax  
    x_max = np.max(x, axis=-1, keepdims=True)  
    exp_x = np.exp(x - x_max) #  
    return exp_x / np.sum(exp_x, axis=-1, keepdims=True)
```

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}} = \frac{e^{x_i-c}}{\sum_j e^{x_j-c}}$$

$$c = \max(x)$$

-